



Caching – Methodology & Strategies

By Vu Van Tiep

Manager, Head Of Web Platform - VNG

Skype: tiepvv

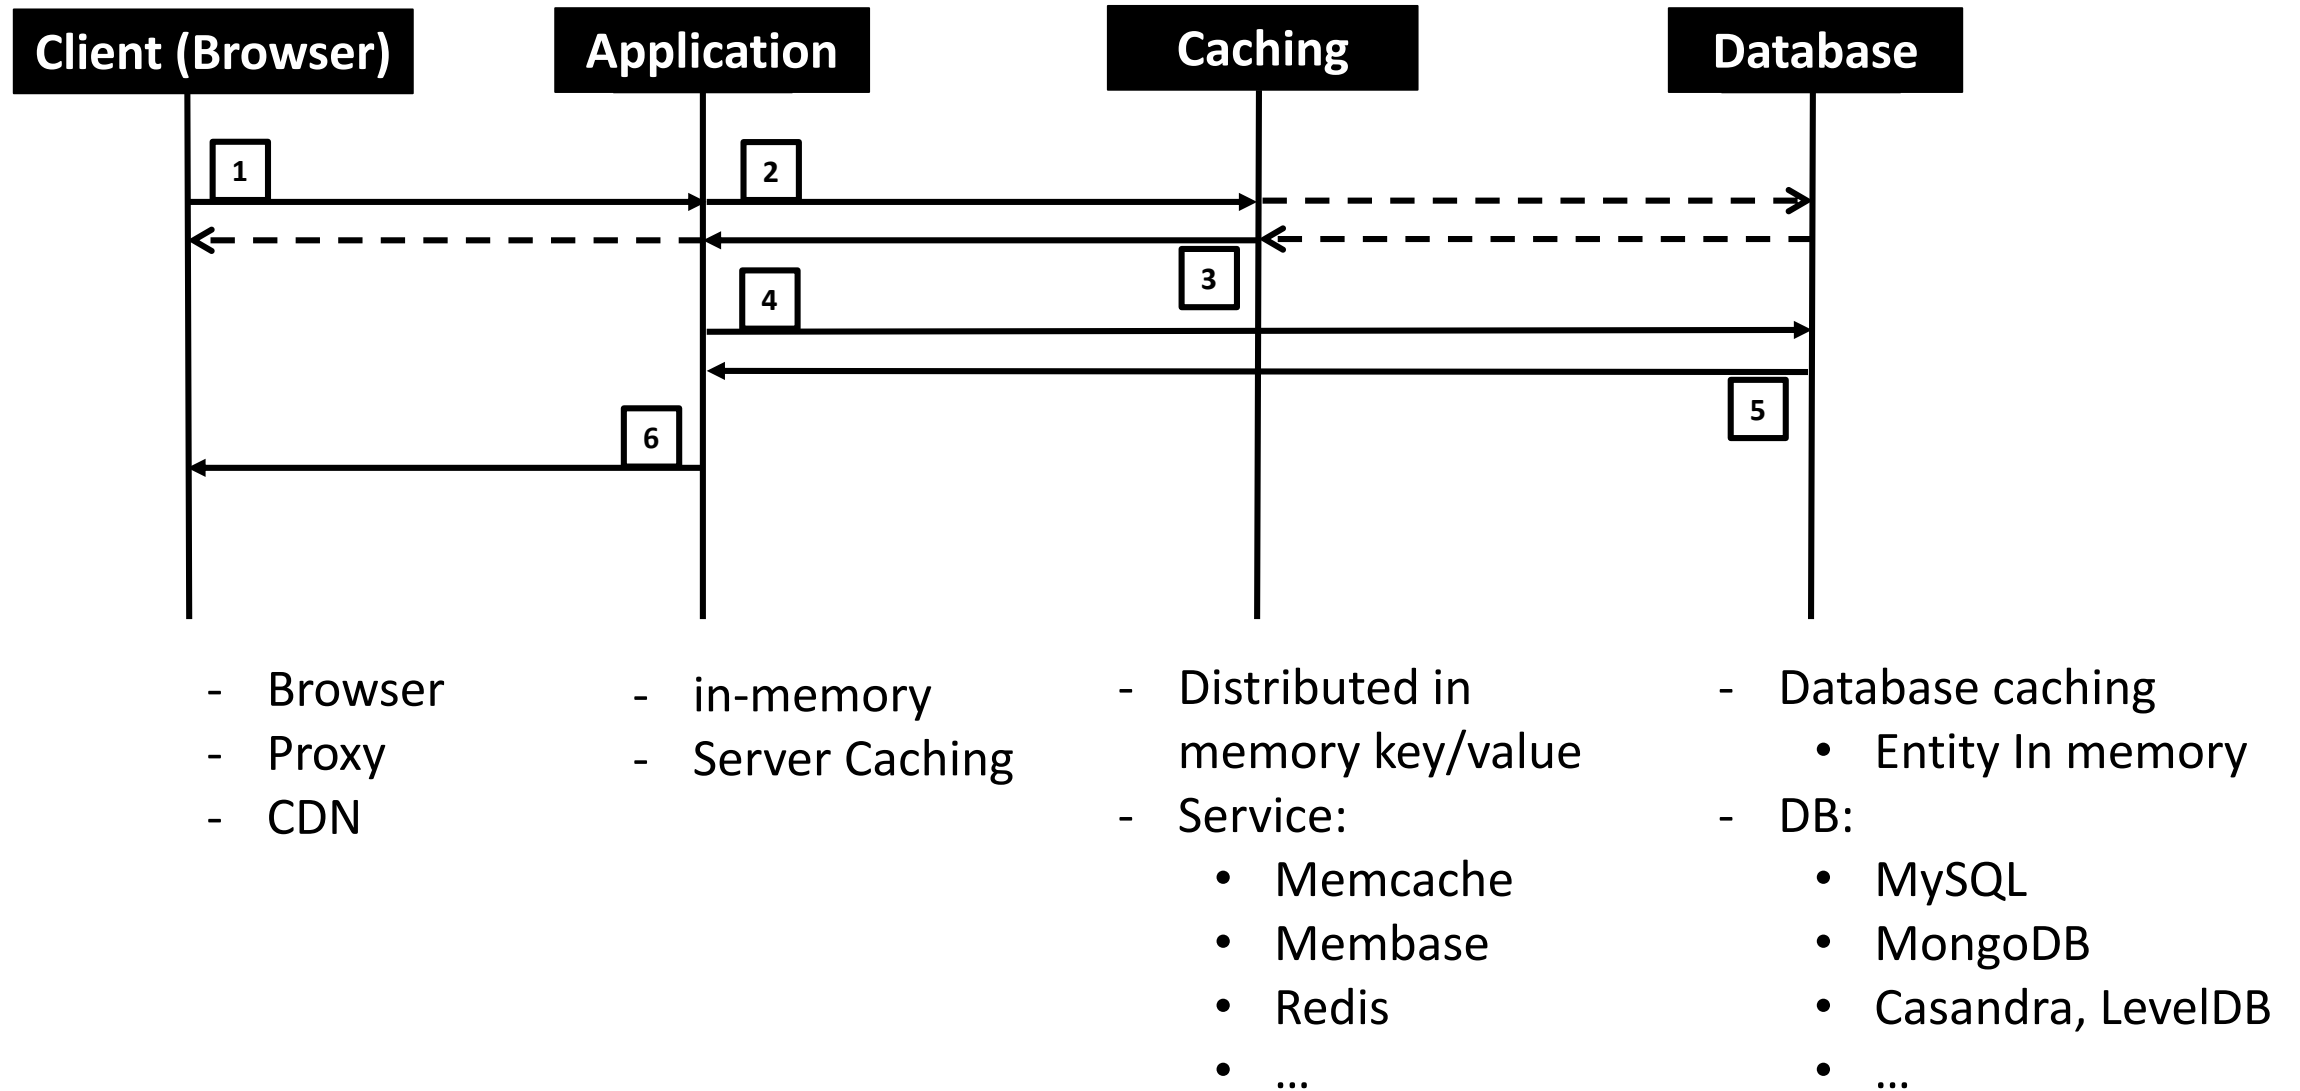
Agenda

- What
- Cache Layer
- Cache Application/Service
 - Methodology
 - Strategies
 - Performance?
- Q&A




Cache – What?

- Is a hardware or software component that stores data so future requests for that data can be **served faster**;
- The data stored in a cache might be the result of an **earlier computation**, or the **duplicate** of data stored elsewhere
- A *cache **hit*** occurs when the requested data can be found in a cache, while a *cache **miss*** occurs when it cannot.
- Cache **hits** are served by reading data from the cache, which is faster than recomputing a result or reading from a slower data store; thus, the more requests can be served from the cache, the faster the system performs.

Cache Layer

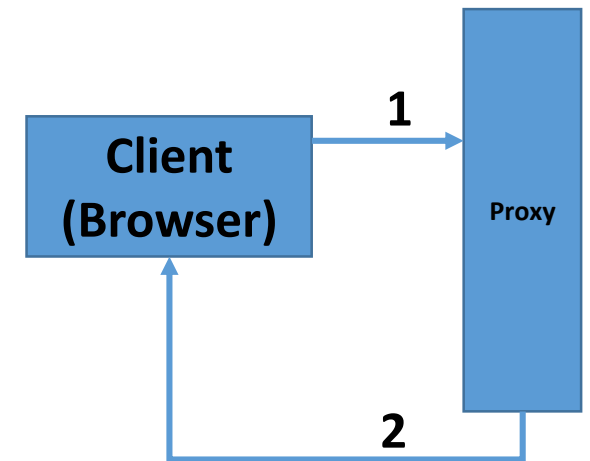


Http Layer Caching – Browser & Proxy



Response Header

- Expires
- Last-Modified
- Etag
- Cache-Control
- Pragma



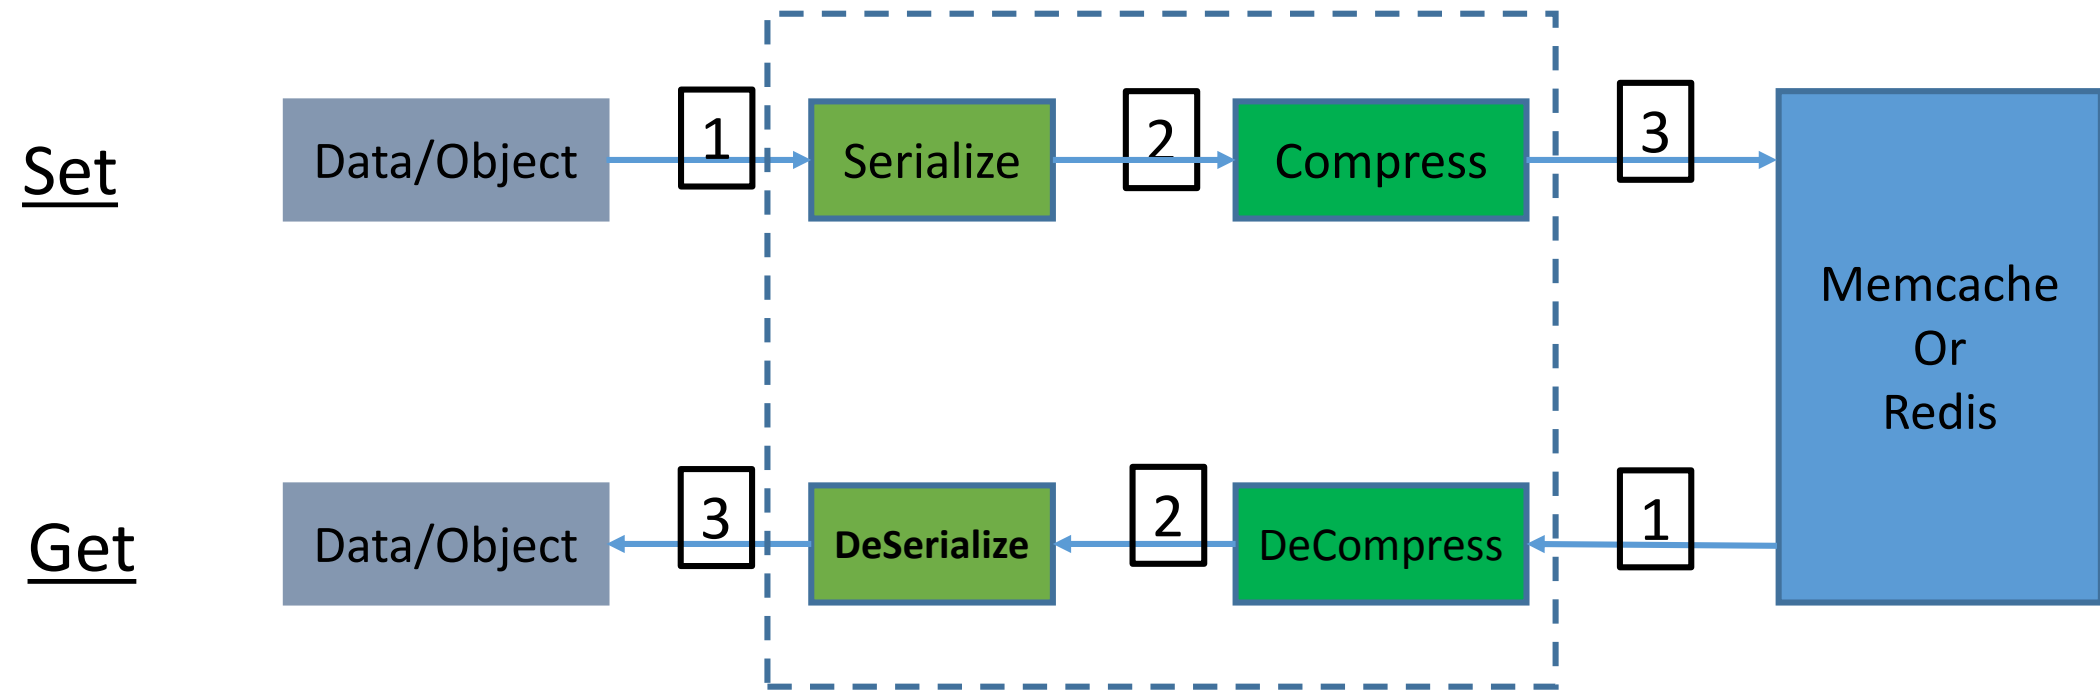
Cache Application/Service

Latency Numbers


L1 cache reference	0.5 ns	
Branch mispredict	5 ns	
L2 cache reference	7 ns	
Mutex lock/unlock	100 ns	
Main memory reference	100 ns	
Compress 1K bytes with Zippy	10,000 ns	
Send 2K bytes over 1 Gbps network	20,000 ns	(1)
Read 1 MB sequentially from memory	250,000 ns	(2)
Round trip within same datacenter	500,000 ns	
Disk seek	10,000,000 ns	1 + 2 < (3)
Read 1 MB sequentially from network	10,000,000 ns	
Read 1 MB sequentially from disk	30,000,000 ns	
Send packet CA->Netherlands->CA	150,000,000 ns	

Note: 1 ns = 10⁻⁹ seconds

Cache: Methodology



Serialize



Performance Comparison

Binary serialize take less space

Text Serialize	Framework	Binary Serialize
<p>Json – size: 167 bytes, se: 6ms, de: 3ms</p> <pre>{ "ProfileInfo": { "userId": 1, "registerDate": 1476200214431, "userName": "tiepvv", "birthDate": 1476200214431, "displayName": "Tiep Vu" }}</pre>	<p>XML, Json ProtoBuf, Thrift Message Pack Apache Arvo ...</p>	<p>Thrift – size: 64 bytes, se: 53ms, de: 1ms</p> <pre>namespace java vietnam.websummit.thrift struct ProfileInfo { 1:i32 userId, 2:string username, 3:string displayName, 4:i32 birthDate, 5:i32 registerDate }</pre>
<p>XML – 235 bytes</p> <pre><?xml version="1.0" encoding="UTF-8" ?> <root> <username>tiepvv</username> <registerDate>1476198686</registerDate> <userId>1</userId> <birthDate>1476198686</birthDate> <displayName>Tiep Vu</displayName> </root></pre>		<p>ProtoBuf – size: 31 bytes, se: 34ms, de: 4ms</p> <pre>syntax = "proto3"; package vietnam.websummit.proto; option java_package = "vietnam.websummit.proto"; message ProfileInfo { int32 userId = 1; string username = 2; string displayName = 3; int32 birthDate = 4; int32 registerDate = 5; }</pre>

Compress/Decompress

Name	Ratio	C.speed	D.speed
		MB/s	MB/s
zstd 0.8.2 -1	2.877	330	940
zlib 1.2.8 -1	2.73	95	360
brrotli 0.4 -0	2.708	320	375
QuickLZ 1.5	2.237	510	605
LZO 2.09	2.106	610	870
LZ4 r131	2.101	620	3100
Snappy 1.1.3	2.091	480	1600
LZF 3.6	2.077	375	790

As a reference, several fast compression algorithms were tested and compared on a Core i7-3930K CPU @ 4.5GHz, using lzbench, an open-source in-memory benchmark by @inikep compiled with GCC 5.4.0, with the Silesia compression corpus.

Caching Strategies

Cache Strategies

1. Optimize Data Size
2. Hot Object
3. In-Process
4. Lazy load
5. Distributed

Optimize Data Size

Serialize	
Method	Size (smaller is better)
Thrift — TCompactProtocol	278 (not bad)
Thrift — TBinaryProtocol	460
Protocol Buffers	250 (winner!)
RMI	905
REST — JSON	559
REST — XML	836

High Level Goals:

- Transparent between multiple programming languages (PHP/Java/Cpp/...)
- Maintain Right balance between:
 - Efficiency (how much time/space?)
 - Availability of existing libraries

Compress			
Name	Ratio	C.speed	D.speed
		MB/s	MB/s
zstd 0.8.2 -1	2.877	330	940
zlib 1.2.8 -1	2.73	95	360
brrotli 0.4 -0	2.708	320	375
QuickLZ 1.5	2.237	510	605
LZO 2.09	2.106	610	870
LZ4 r131	2.101	620	3100
Snappy 1.1.3	2.091	480	1600
LZF 3.6	2.077	375	790

Hot Object => Eviction Algorithms



Data



Server

How to cache?



Eviction Algorithms: Least Frequently Used (LFU)

A	B	C	D	E	F
7	6	3	2	4	1

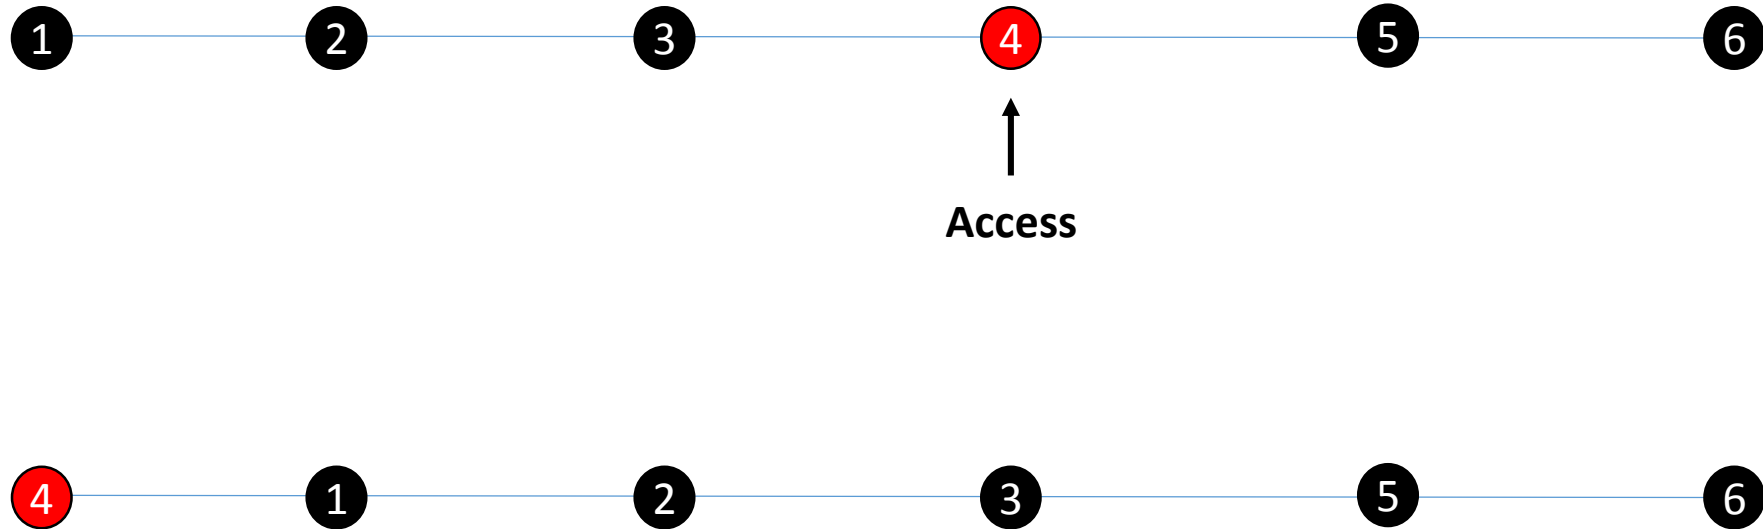
↑
Access

A	B	C	D	E	F
7	6	3	3	4	1

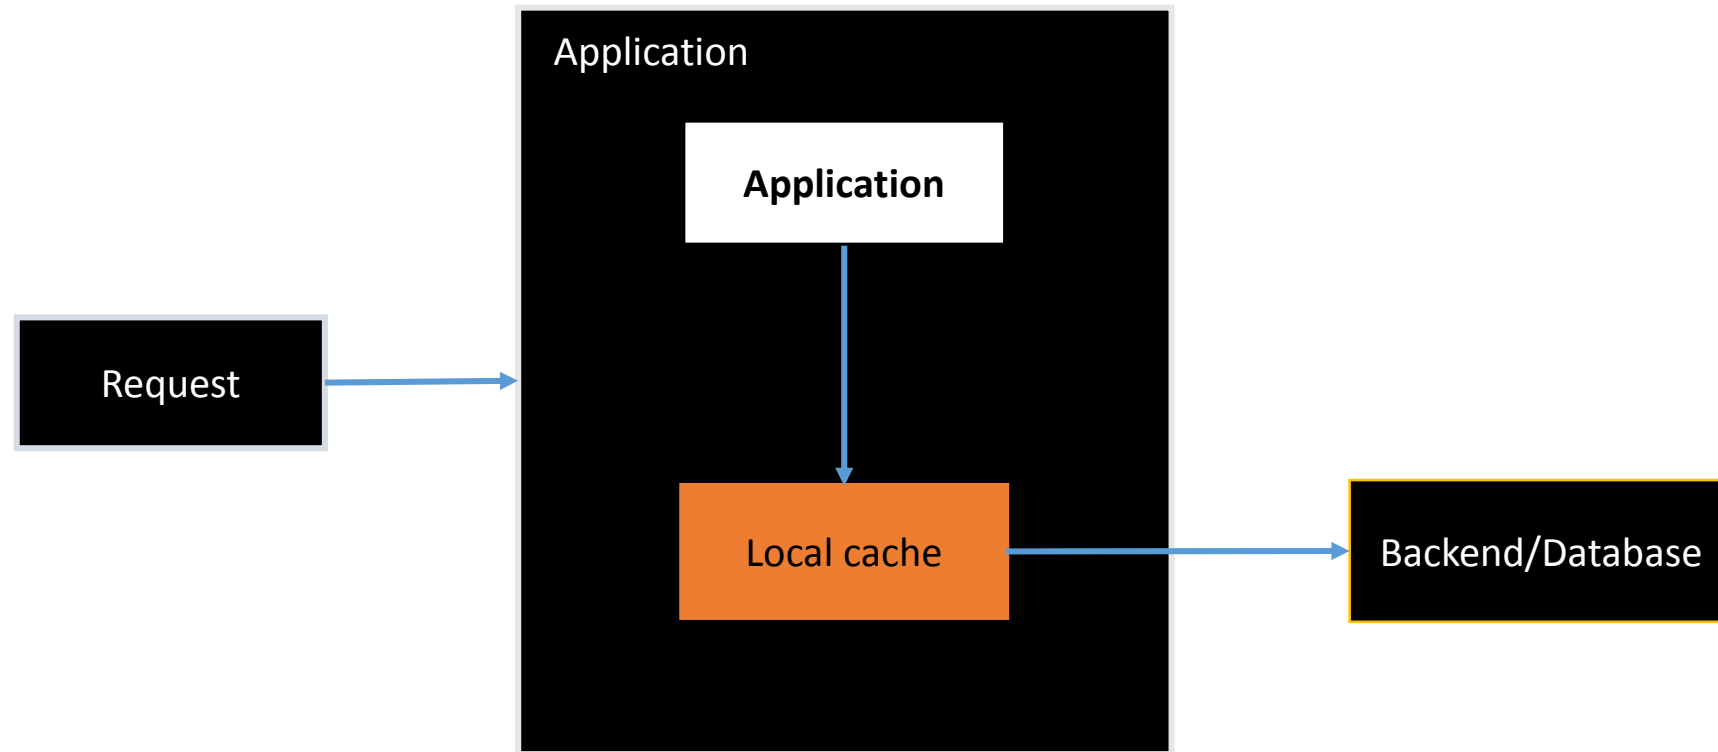
↑
Most Frequently
Used

↑
Least Frequently
Used


Eviction Algorithms: Least Recently Used (LRU)



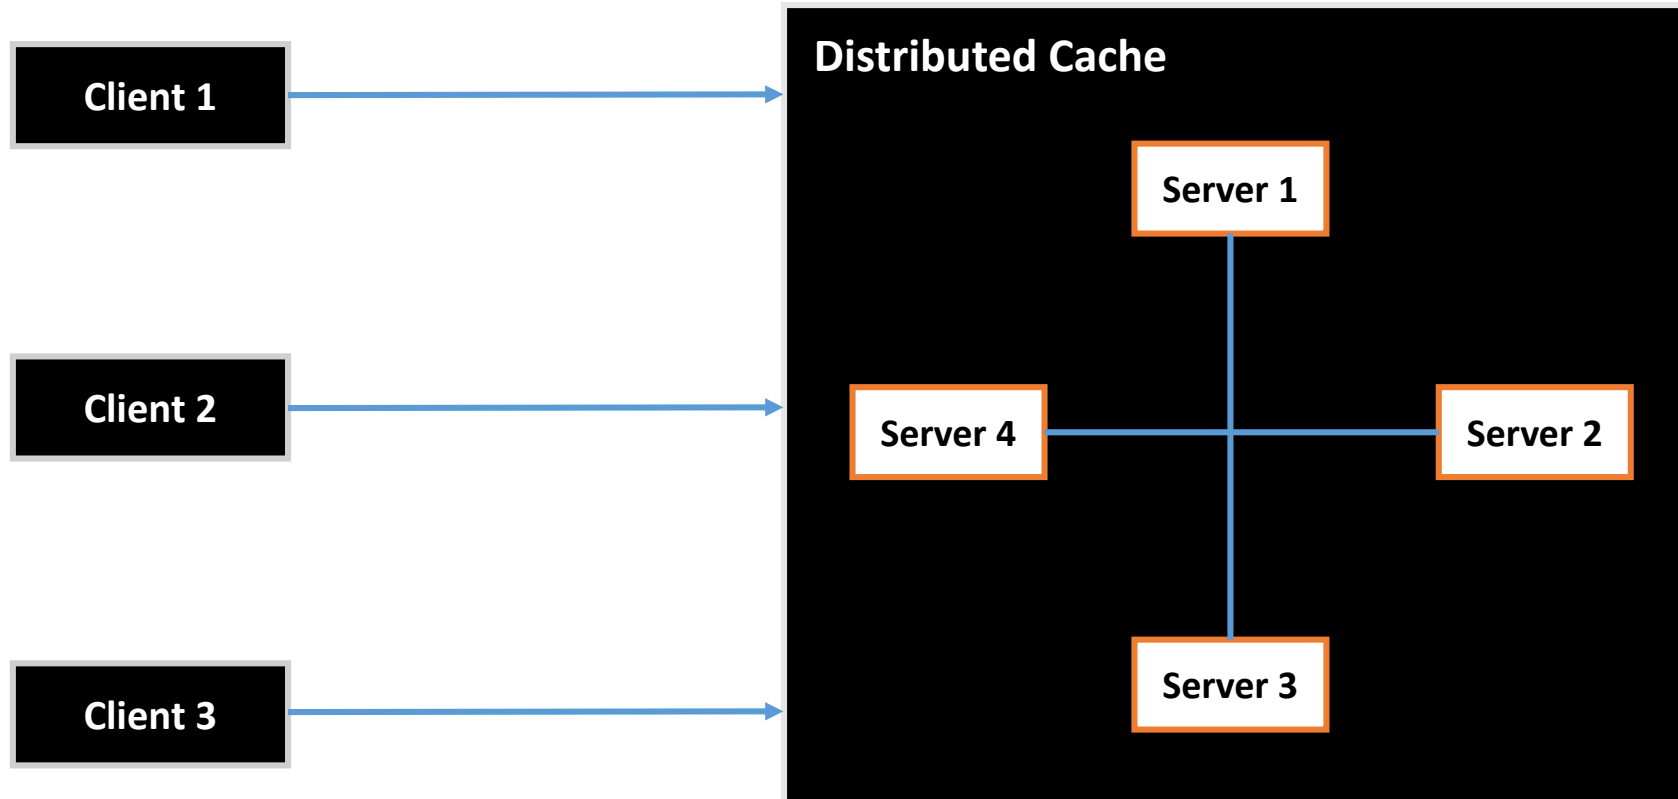
In-Process



Lazy Load




Distributed Cache



Performance: Profiler is the key

Method	Total Request	Last RPS/ Last Exe Time	Max RPS	Min RPS	Avg RPS	Max ExeT	Min ExeT	Avg ExeT
get	2929261981	21,7,14,18,20 0,0,0,0,0	61	2	21.3999996	0	0	0.01778333
multiGet	0	0,0,0,0,0 0,0,0,0,0	0	0	0	0	0	0
enable	7844039	143,132,132,209,165 0,0,0,0,0	231	44	140.983337	0	0	0.19936667
disable	0	0,0,0,0,0 0,0,0,0,0	0	0	0	0	0	0
remove	7844039	143,132,132,209,165 0,0,0,0,0	231	44	140.983337	0	0	0.21293333
getStat	6690	1,1,1,1,0 0,0,0,22,0	1	0	0.0666666701	22	0	0.41006669
Cache: item	Cache hit(%)		Cache miss(%)					
23982034	99.1956482		0.804351926					

Q&A



Thank You!

Contact info:

- Name: Vu Van Tiep
- Email: tiepvv@vng.com.vn / vuvantiep@gmail.com
- Skype: tiepvv



Reference

- [https://en.wikipedia.org/wiki/Cache \(computing\)](https://en.wikipedia.org/wiki/Cache_(computing))
- <http://www.slideshare.net/IgorAnishchenko/pb-vs-thrift-vs-avro>
- <https://github.com/facebook/zstd>
- <http://www.codeproject.com/Articles/56138/Consistent-hashing>